


A single screenshot from the middle of the animation sequence in a New York Times visualization depicting income inequality:

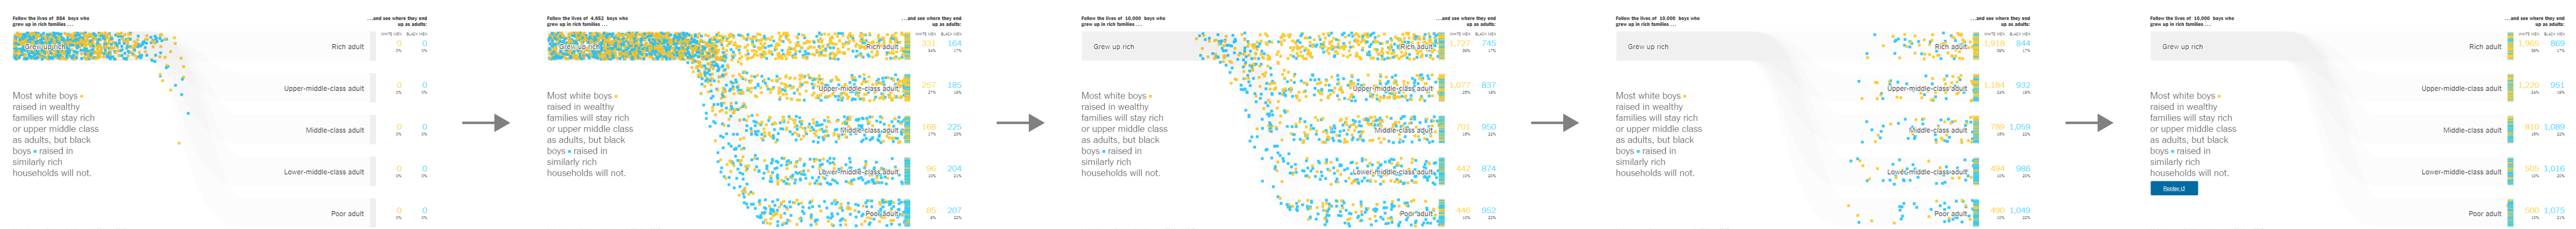
**Follow the lives of 10,000 boys who grew up in rich families ...**

...and see where they end  
up as adults:



As the animation progresses, dots move from the left side:  $P(\text{race} \mid \text{grew up rich})...$

... to the right side:  
 $P(\text{race} \mid \text{income class as an adult} \wedge \text{grew up rich})$



Note the final conditional distribution on the right-hand side is **not**  $P(\text{income class as an adult} \mid \text{race} \wedge \text{grew up rich})$ , which is the distribution one should depict if one is interested in what income class a black or white boy from a rich family will grow up to be in (and is implicitly the distribution referred to in the text annotation on the visualization, even if it is not depicted).